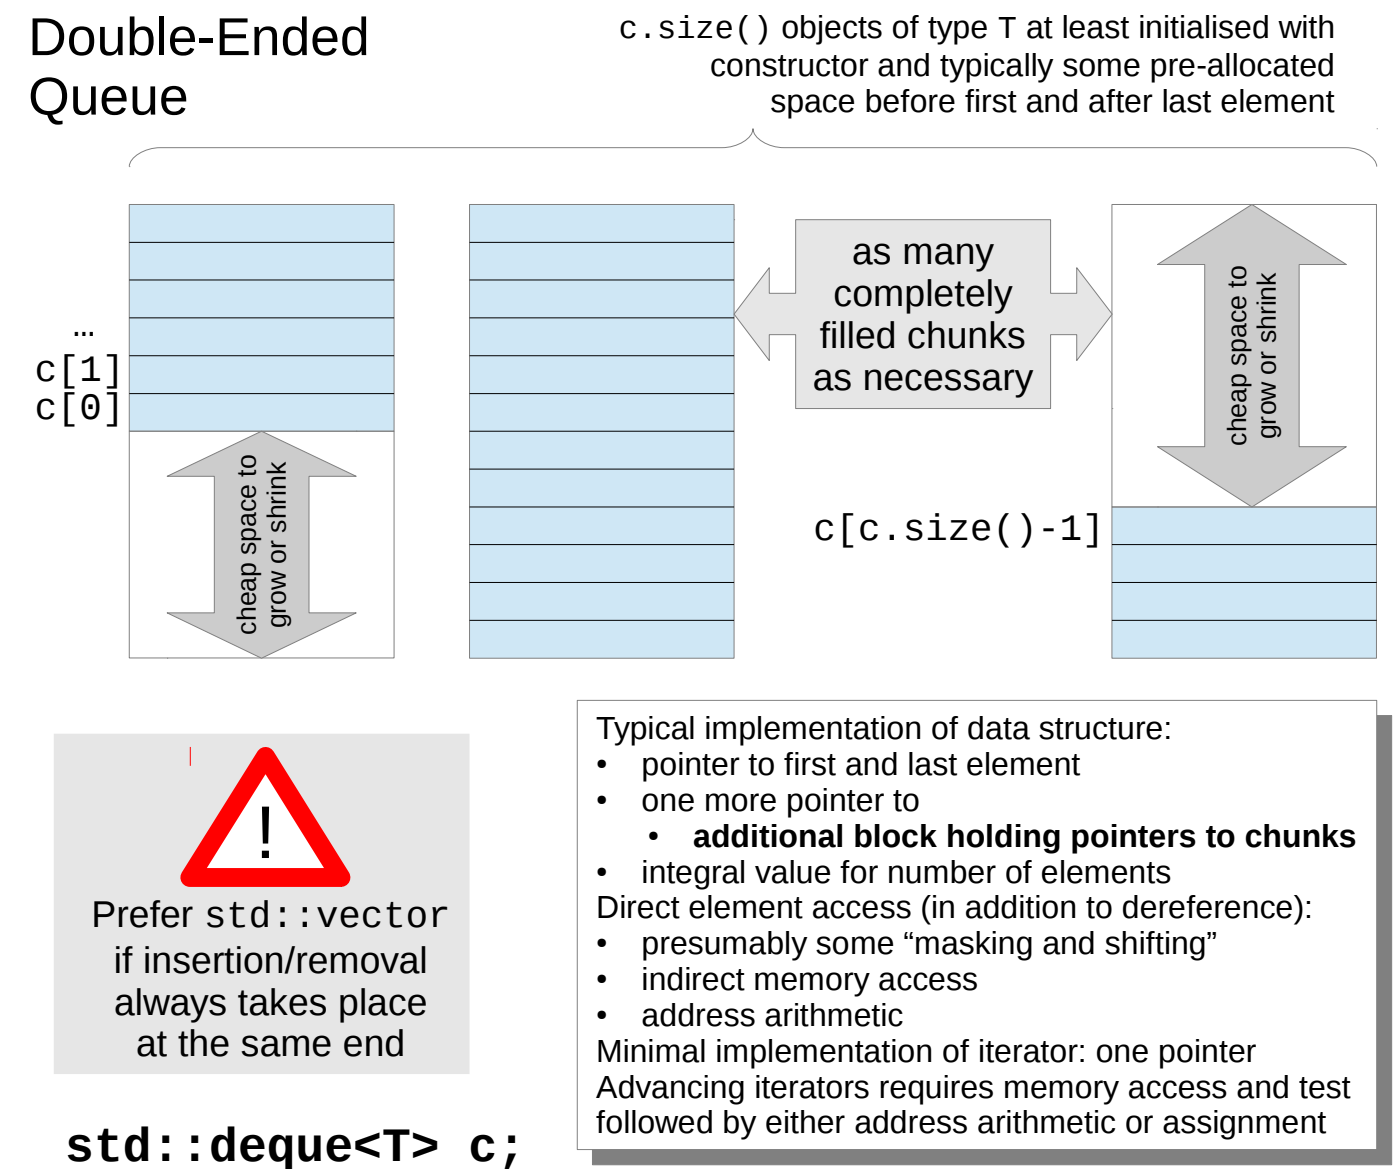



Double-Ended Queue



`std::list<T> c;`

`c.size()` objects of type T at least initialised with constructor



Typical implementation:

- two pointers per element**
- pointer to first and last element
- integral value for number of elements

Direct element access not supported!

Minimal implementation of iterator: one pointer


Advancing iterators requires memory access followed by assignment

Double Linked List

Substantial overhead if `sizeof(T)` is small.

`std::forward_list<T> c;`

objects of type T initialised with constructor



Singly Linked List

Use `c.empty()` to check whether elements exist.

Typical implementation:

- one pointer per element**
- only pointer to first element
- number of elements not stored!**

Direct element access not supported!

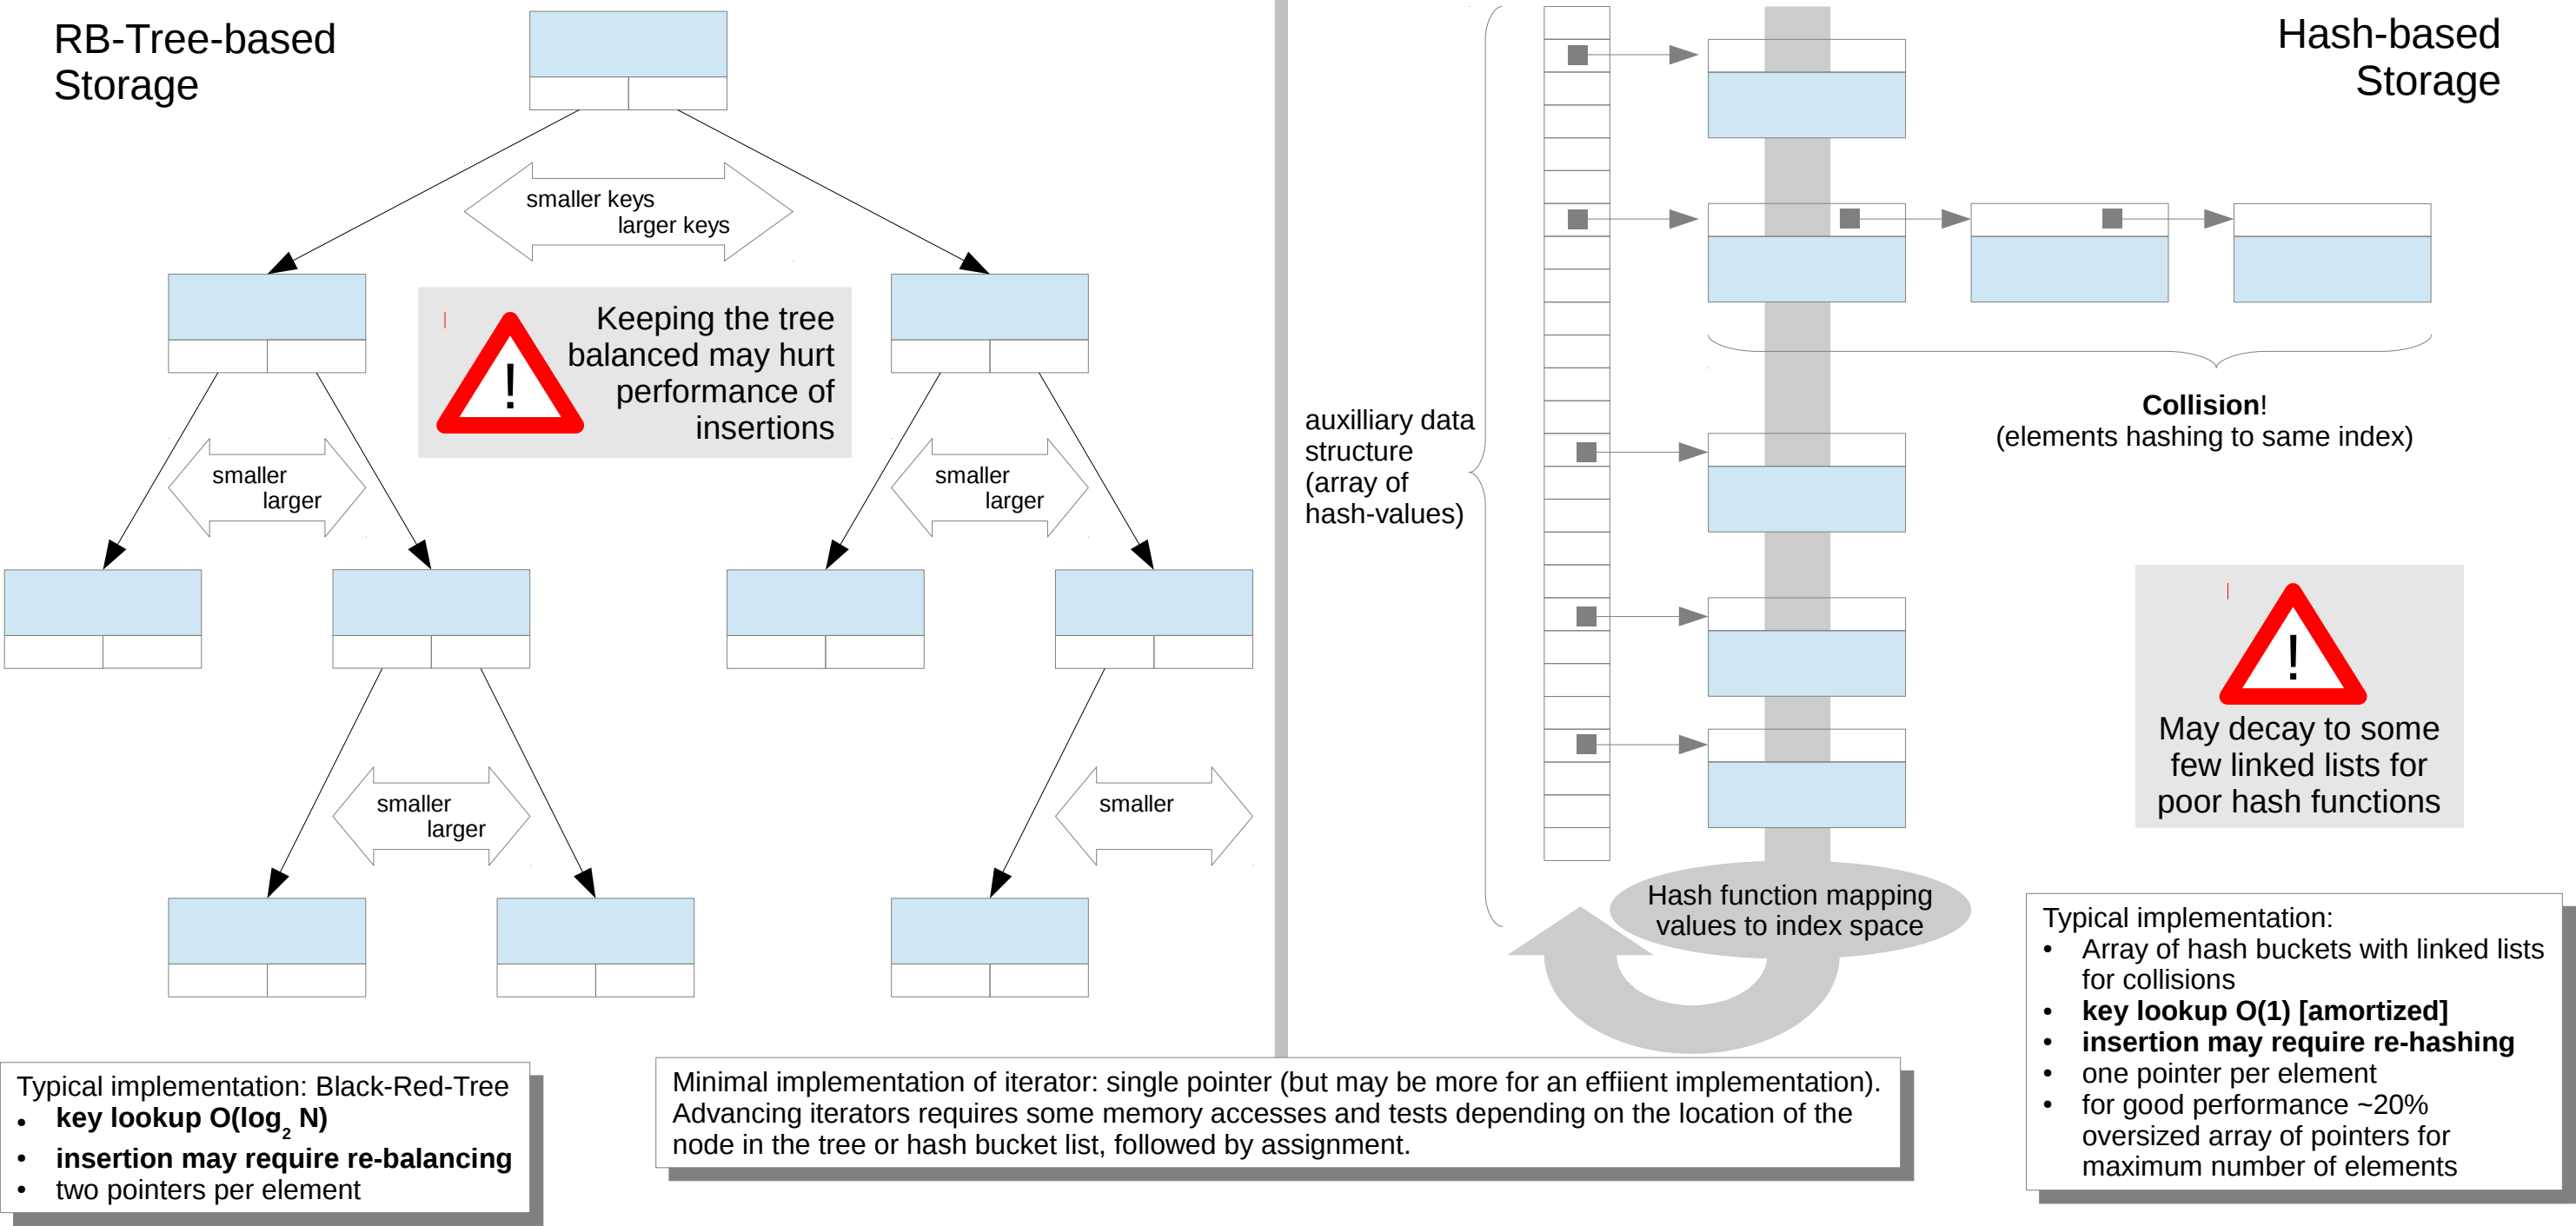
Minimal implementation of iterator: one pointer

Advancing iterators requires memory access followed by assignment

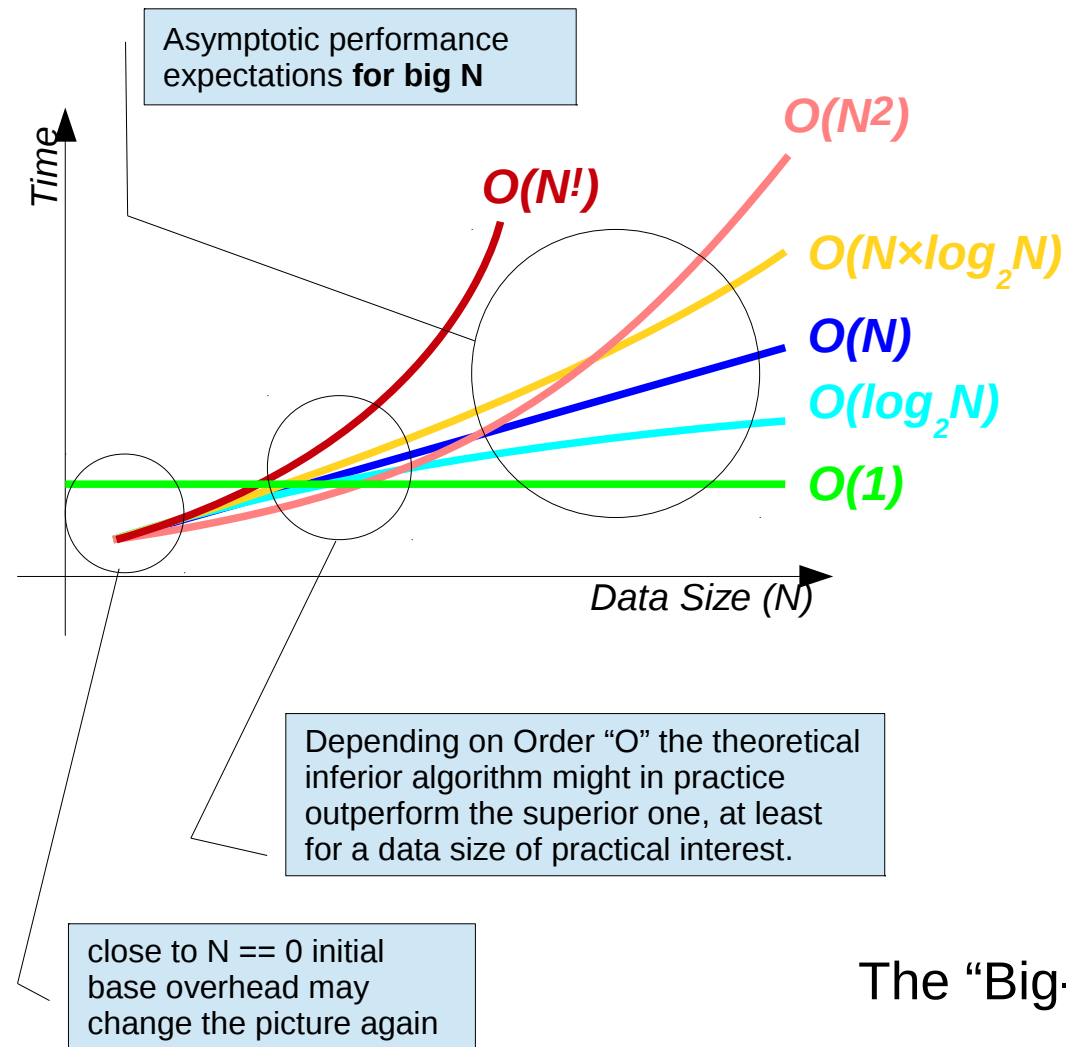
STL – Sequence Container Classes

(cc) BY-SA: Technische Beratung für EDV, Dipl.-Ing. Martin Weitzel, Germany, <http://tbfe.de>

Contained elements	STL Class Name		Restrictions
objects of type T	<code>std::set</code>	<code>std::unordered_set</code>	unique elements guaranteed
	<code>std::multiset</code>	<code>std::unordered_multiset</code>	multiple elements possible (comparing equal to each other)
pairs of objects of type T_1 (key) and type T_2 (associated value)	<code>std::map</code>	<code>std::unordered_map</code>	unique keys guaranteed
	<code>std::multimap</code>	<code>std::unordered_multimap</code>	multiple keys possible (comparing equal to each other)



STL – Associative Container Classes




The "Big-O"-Notation

Practical Consequences of "Big-O"

N =	1	10	100	1000	1e6 one Million	1e9 one Billion	1e12 one Billion
$O(1)$	1sec	1sec	1sec		1sec	1sec	1sec
$O(\log_2(N))$	0,5sec	0,63sec	0,8sec		~2sec	~4sec	~8sec
$O(N)$	1msec	10msec	0,1sec	1sec	~¼ hour	~2½ weeks	~30 years
$O(N \times \log_2(N))$	0,1msec	6msec	75msec		~½ hour	~2 months	~300 years
$O(N^2)$	1µsec	0,1msec	10msec		~2½ weeks	~Cro-Magnon to modern	~20 × age of universe

Complexity	(commonly known as)	Example
$O(1)$	constant time	lookup in hash-based data structure
$O(\log_2(N))$	logarithmic time	lookup in tree-based data structure
$O(N)$	linear time	lookup in sequential data structure
$O(N \times \log_2(N))$	linear-logarithmic time	quick sort
$O(N^2)$	quadratic time	bubble sort



... and
Amortized
Analysis

Some STL Quirks


`std::binary_search` is specified as **$O(\log_2 N)$** even though it doesn't require random access iterators.

`std::vector` specifies its `push_back` operation as **amortized $O(1)$** even though there typically are repeated expensive copy operations on always larger memory areas;

Add item to <i>linear list</i> vs. <i>binary tree</i> :	$O(1)$	$O(\log_2(N))$
Add N elements to reserved memory, if filled up copy all existing data with duplicating reservation vs. extending reservation by constant amount:	amortized $O(1)$	amortized $O(N^2)$
Linear search vs. binary search:	$O(N)$	$O(\log_2(N))$
Optimal sort vs. bubble sort:	$O(N \times \log_2(N))$	$O(N^2)$
For N different values create all rotations vs. all permutations:	$O(N)$	$O(N!)$
Shortest of N paths vs. travelling sales-man problem:	$O(N)$	$O(N^2 \times \log_2(N))$



Algorithmic Complexity

Effects of Alternative Approaches






Emphasizing Element Access:

- Iterator points **onto** elements
- **must not be dereferenced in end position!**



Forward Iterator



Reverse Iterator





Iterator for Empty Container

Order defined by

- **insertion** (deletion, explicit sorting ...) for vector, deque, list, forward_list
- **element order** for set and multiset
- **key order** for map and multimap
- implementation for unordered_-containers(i.e. technically unspecified)

(front) container filled with some elements (back)



Increment operation moves iterator by one element

Iterator dereferencing accesses

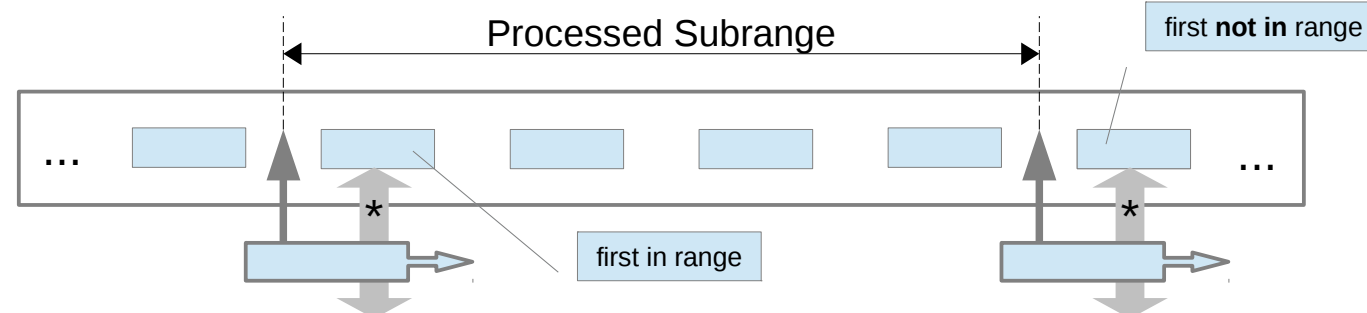
- element value for most containers ...
- ... **except** for all kinds of map-s where a struct with elements first (key) and second (value) is returned

Essentials




Full vs. ...

... Partial Container Processing




STL – Container Iterators

(cc) BY-SA: Technische Beratung für EDV, Dipl.-Ing. Martin Weitzel, Germany, <http://tbfe.de>





Emphasizing Current Position:

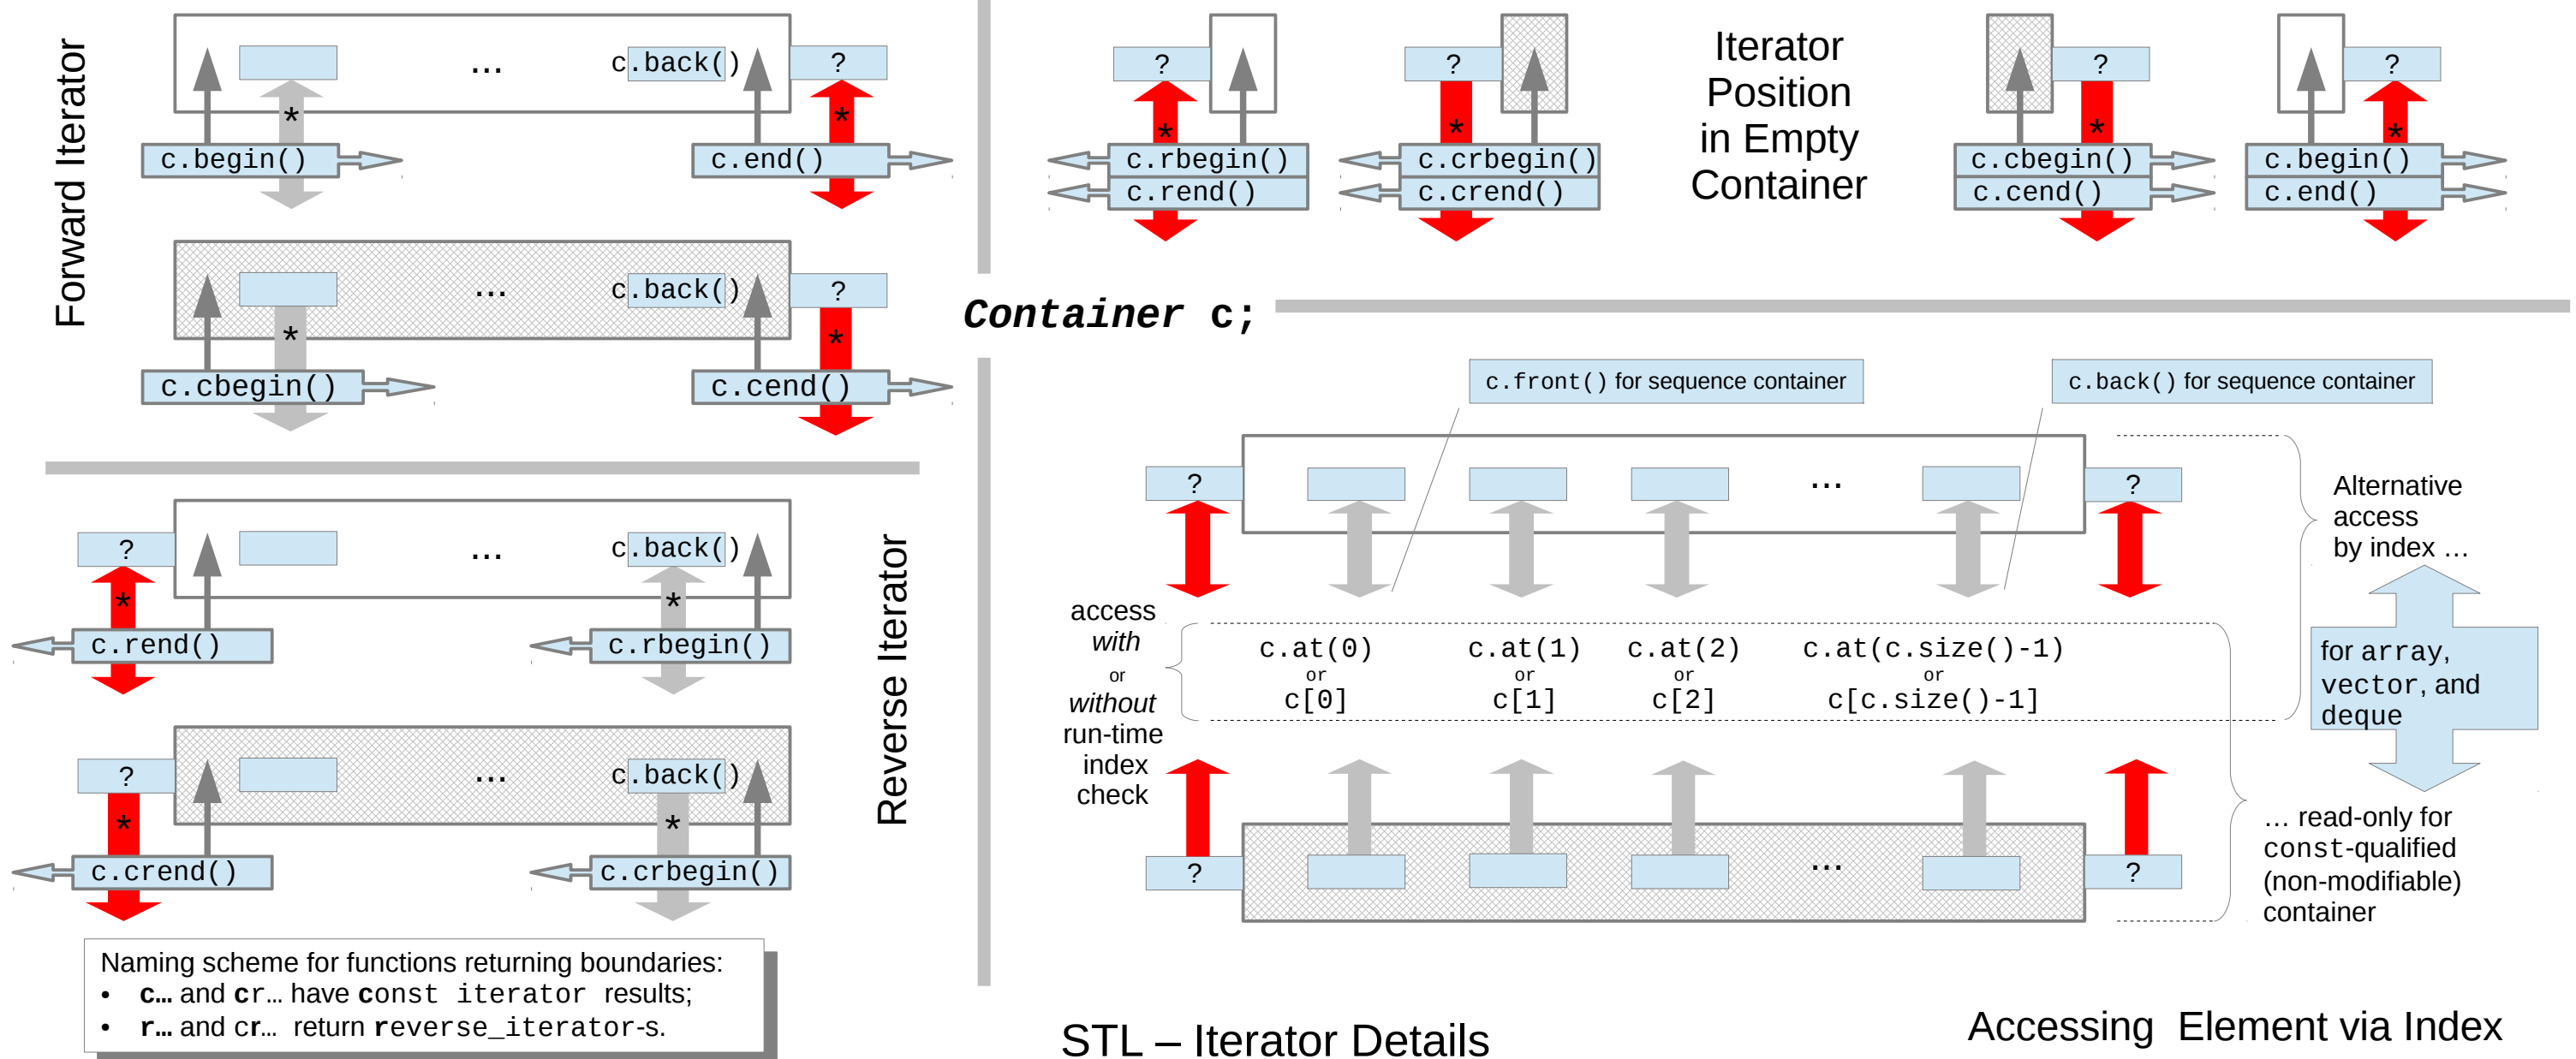
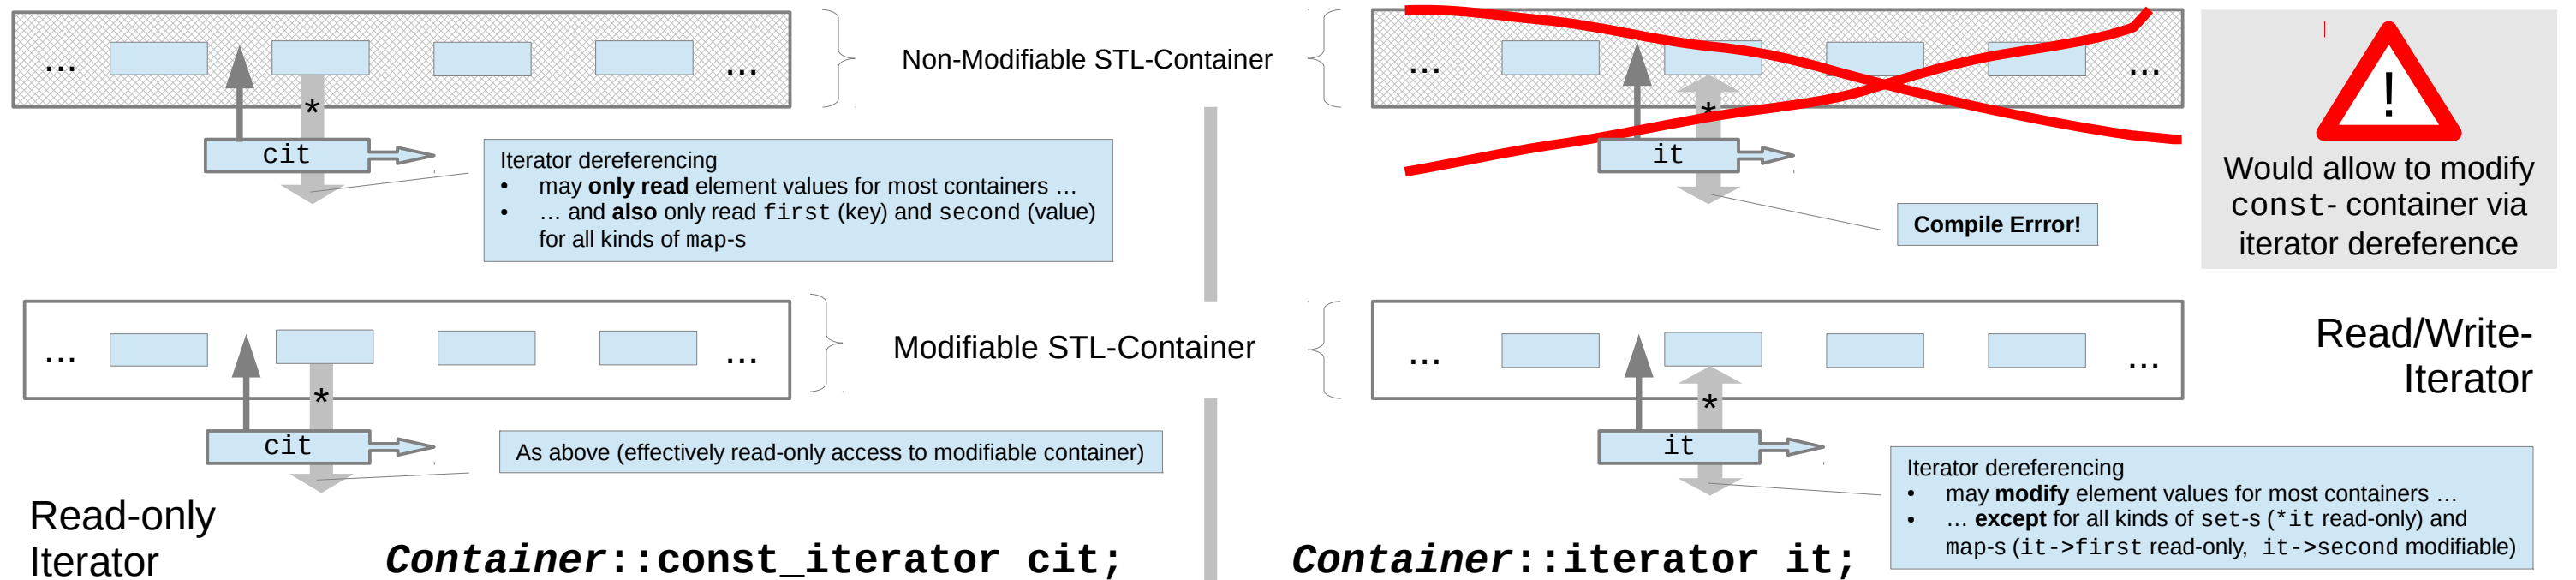
- Iterator points **between** elements
- **accessed element lies in direction of move**




Forward Iterator



Reverse Iterator






Operations of Unidirectional Iterators

	Effect	Remarks
<i>*it</i>	access referenced element	undefined at container end
<i>++it</i> <i>it++</i>	advance to next element (usual semantic for pre-/postfix version)	
<i>it == it</i>	compare for identical position	operands must denote existing element or end of same container
<i>it != it</i>	compare for different position	

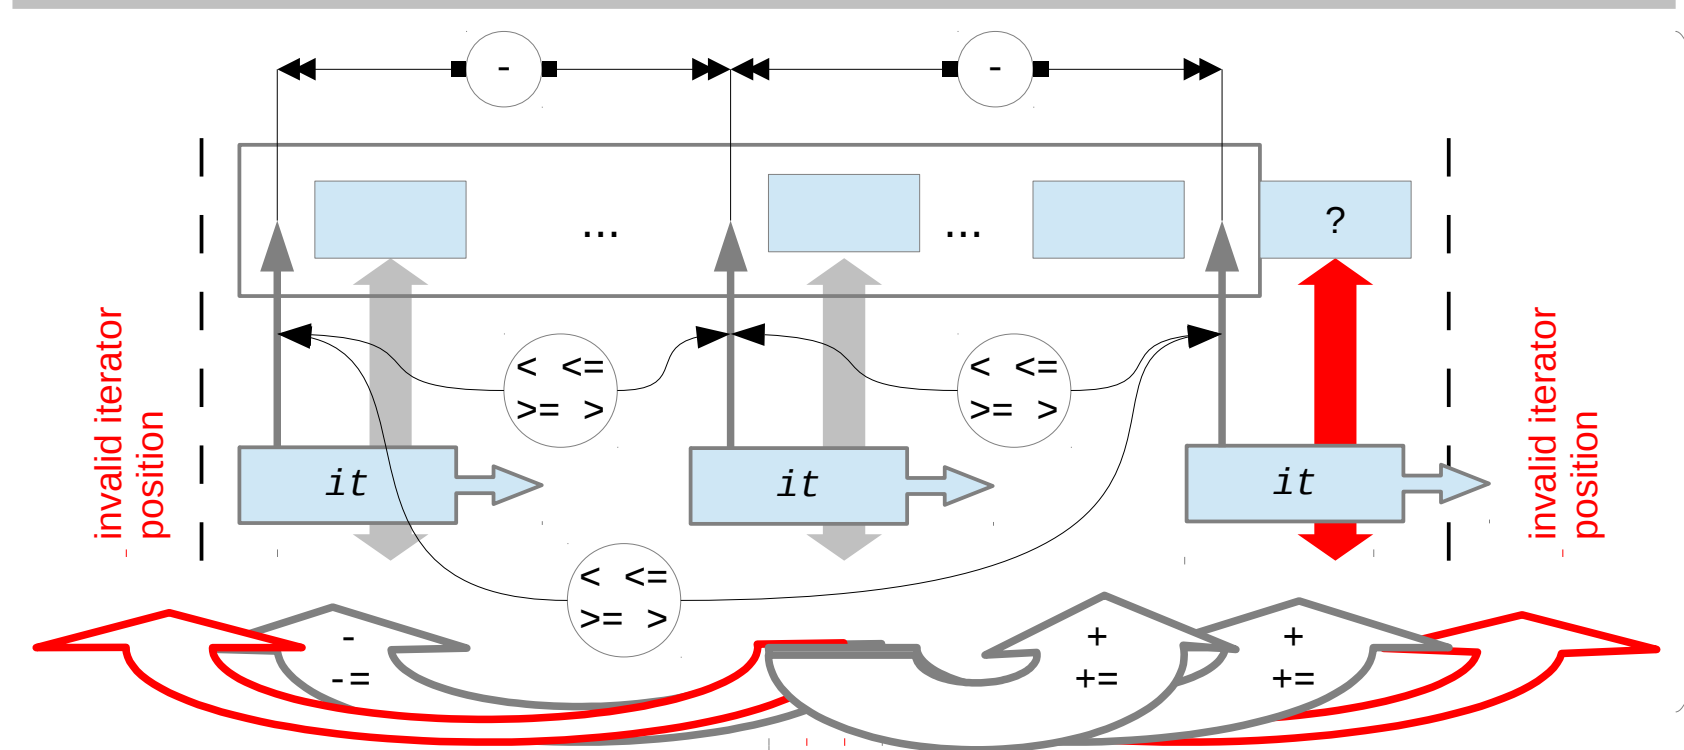
Additional Operations of Bidirectional Iterators

	Effect	Remarks
<i>--it</i> <i>it--</i>	advance to previous element (usual semantic for pre-/postfix version)	undefined at container begin



Additional Operations of Random Access Iterators

	Effect	Remarks
<i>it + n</i> <i>it += n</i>	<i>it</i> advanced to <i>n</i> -th next element (previous if <i>n</i> < 0)	resulting iterator position must be inside container (denote existing element or end)
<i>it - n</i> <i>it -= n</i>	<i>it</i> advanced to <i>n</i> -th previous element (next if <i>n</i> < 0)	
<i>it - it</i>	number of increments to reach rhs <i>it</i> from lhs <i>it</i>	operands must denote existing element or end of same container
<i>it < it</i>	true lhs <i>it</i> <u>before</u> rhs <i>it</i>	
<i>it <= it</i>	true if lhs <i>it</i> <u>not after</u> rhs <i>it</i>	
<i>it >= it</i>	true if lhs <i>it</i> <u>not before</u> rhs <i>it</i>	
<i>it > it</i>	true if lhs <i>it</i> <u>after</u> rhs <i>it</i>	



STL – Iterator Categories

	Container Dimension													
Library	STL									Standard Strings	Iterator Interface to I/O-Streams		e.g. Boost	Others
Kind of Container	Sequential Containers				Associative Containers									
Data Structure	Random Access			Sequential Access		Tree	Hash	Tree	Hash		I/O operations for some type T			
Class Name	array	vector	deque	list	forward_list	set	unordered_set	map	unordered_map	string wstring ...	istream_iterator	ostream_iterator		
						multiset	unordered_multiset	multimap	unordered_multimap					
Iterator Category	Random Access Iterators			Bidirectional Iterators	Unidirectional Iterators	Bidirectional Iterators				Random Access Iterators	Input Iterators	Output Iterators		
Dereferenced Iterator	accesses element							accesses key-value-pair		single character	single item of type T			
	operations available via iterators													

Algorithm Dimension	STL	Access: • find • search • ...
		Modify: • remove • sort • ...
		Miscellaneous: • count • mimmax_seq • ...
	e.g. Boost	Algorithm: • join • ... String_algo: • trimleft • trimright • ... • ...
Others		

operations expected from iterators




Failure to comply will cause a compile-time error, typically with respect to the header file that defines the algorithm.

Iterators as "Glue" to connect Containers with Algorithms



Failure to comply will either cause a compile-time error or show at runtime and may depend on the kind of container.

elements still physically present though no longer logically part of the container



"Removing" Elements ... returns "New End"

Use of iterators to specify container elements to process:

- starting point is the first element to process
- ending point is the first element **not** to process
- whole container is specified via its begin() and end()




always valid for dereferencing

... Return Success ...

not necessarily valid for dereferencing!

... or Failure



Input Iterators Semantic Restrictions

- * must only be used for read access
- ++ must follow each read exactly once

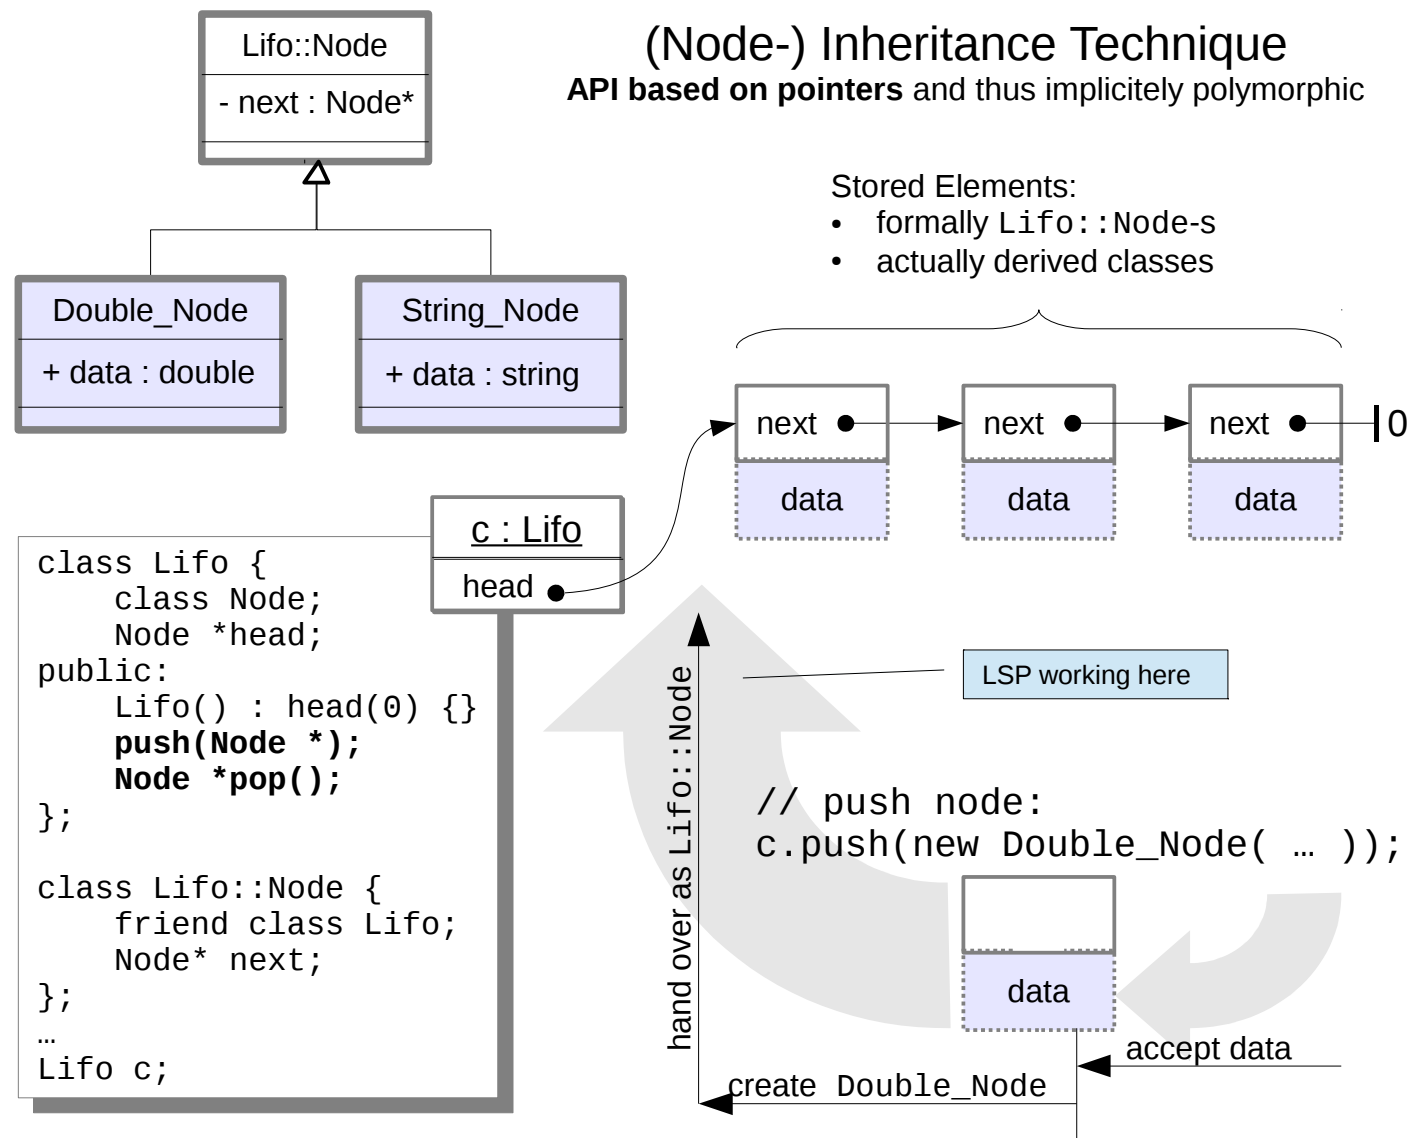
Output Iterators Semantic Restrictions

- * must only be used for write access
- ++ must follow each write exactly once

STL – Iterator Usages

(Node-) Inheritance Technique

API based on pointers and thus implicitly polymorphic



```

// pop node (double expected):
if (auto *p = dynamic_cast<Double_Node *>(c.pop())) {
    // process node data:
    ... p->data ...
    ...
    // owning Node now!
    delete p;
}
    
```

unexpected node types may cause memory leak with this coding style!

dynamic down-cast may return nullptr

extend to `Double_Node*`

get `Lifo::Node*`

```

// zero run-time overhead (may cause undefined behavior):
... static_cast<Double_Node *>(p)->data ...
    
```

```


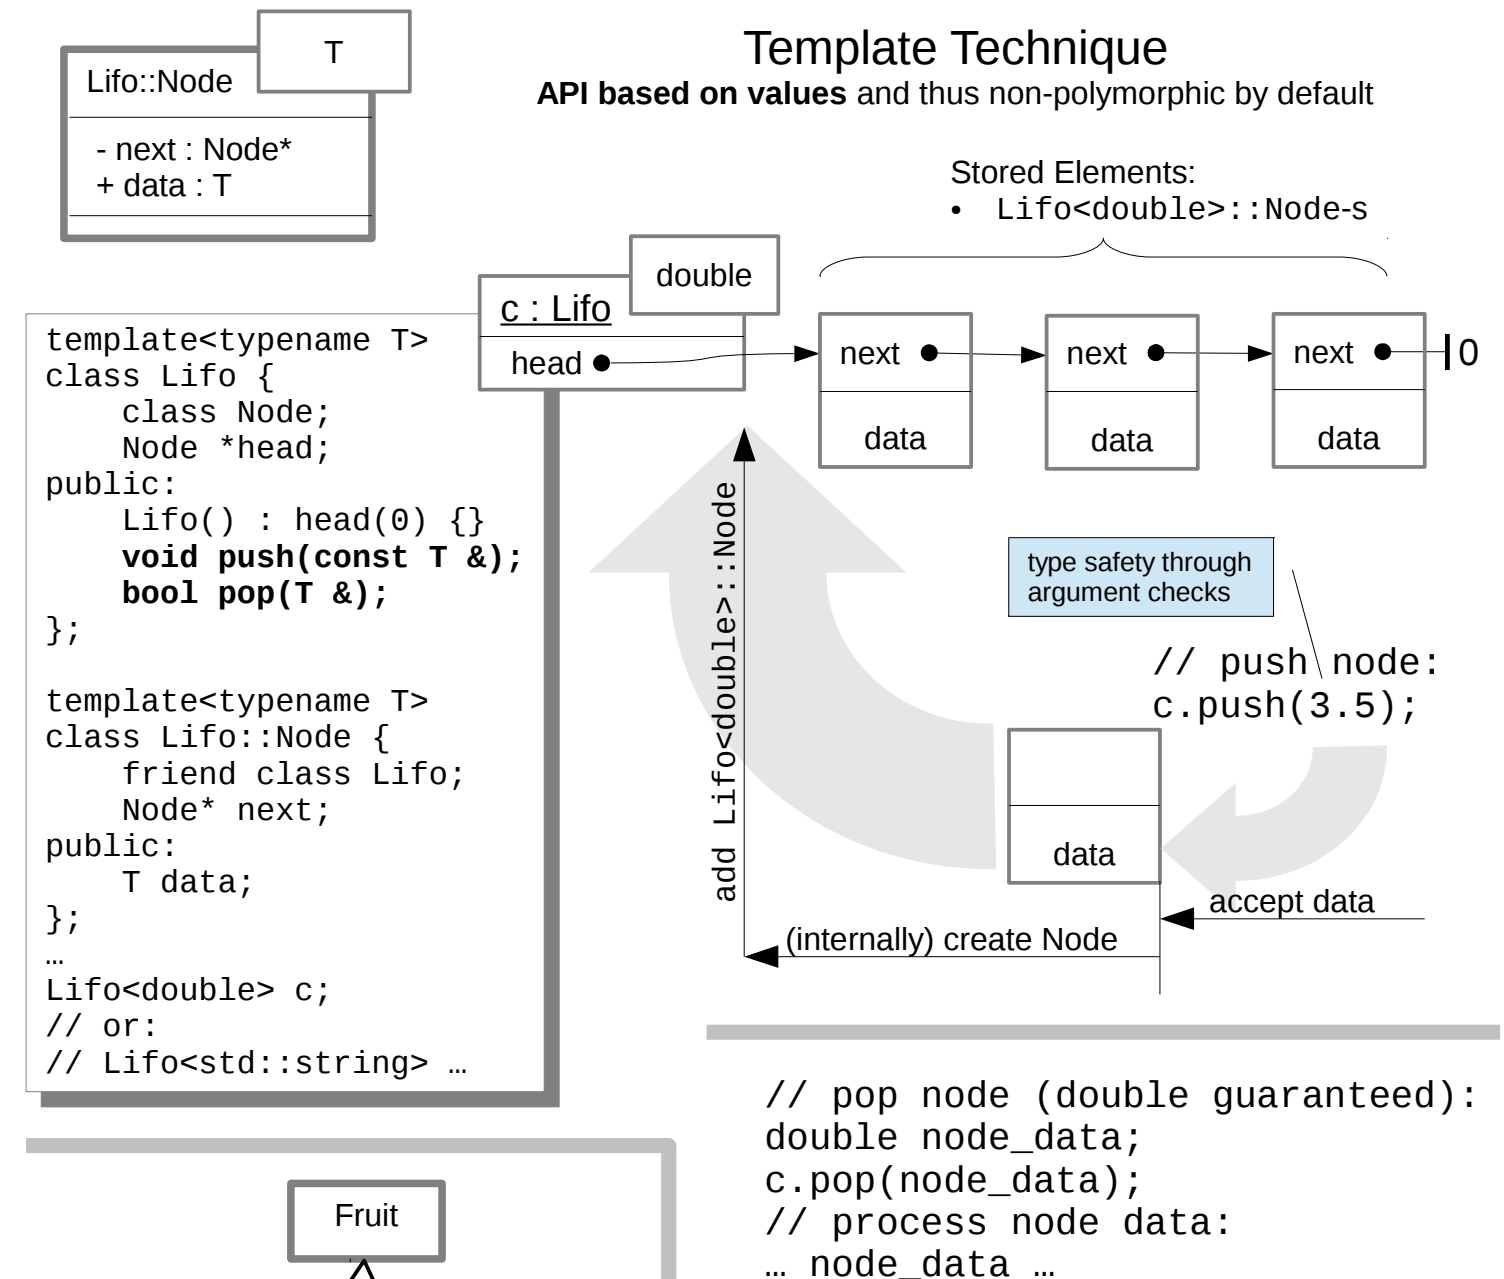
// safe short-hand (may throw):
... dynamic_cast<Double_Node &>(*p).data ...
    
```

Container Implementation Techniques

(cc) BY-SA: Technische Beratung für EDV, Dipl.-Ing. Martin Weitzel, Germany, <http://tbfe.de>

Template Technique

API based on values and thus non-polymorphic by default



Polymorphic Elements

```

Fifo<Fruit*> basket;
...
basket.push(new Apple( ... ));
...
fruit *f;
basket.pop(f);
    
```

for polymorphic elements containers must use pointers explicitly

now points to an Apple

Non-Polymorphic Elements

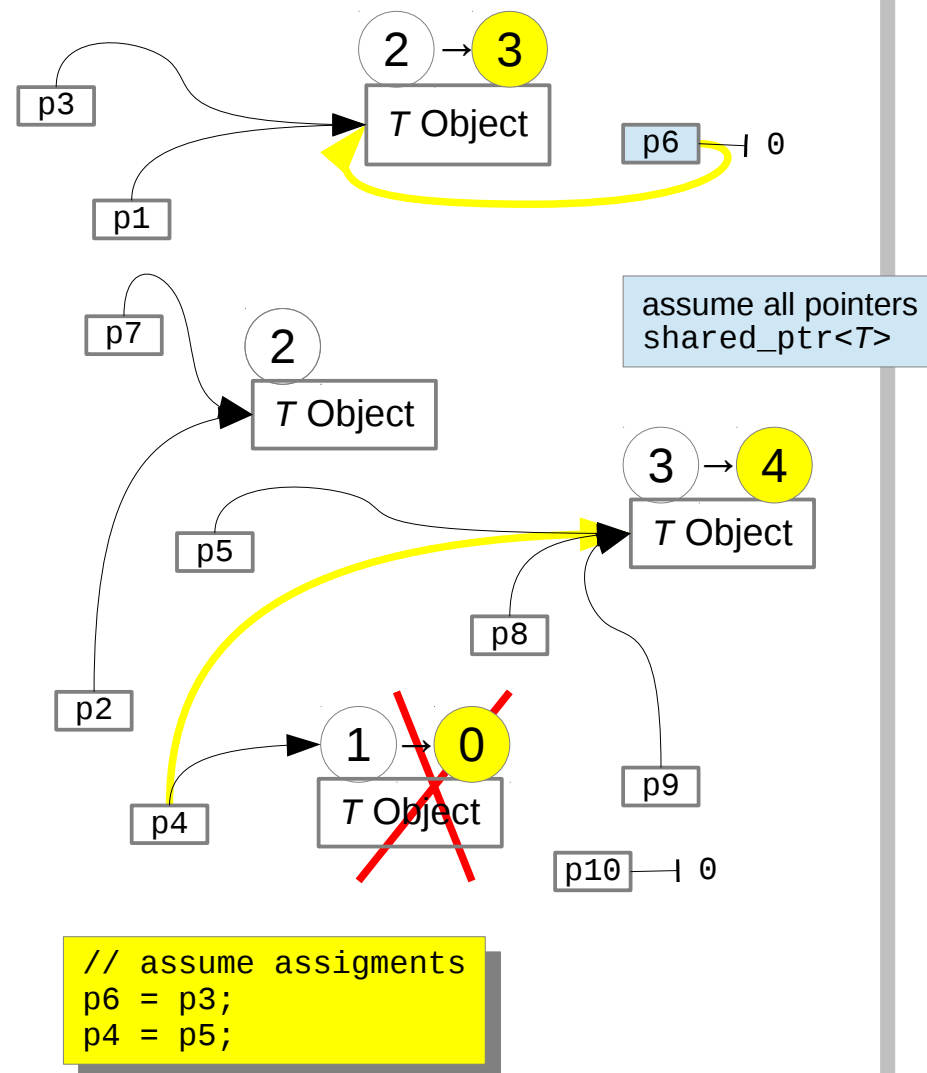
```

Fifo<Fruit> basket;
Apple a;
Banana b;
...
basket.push(a);
basket.push(b);
...
Fruit f;
basket.pop(f);
...
Basket.pop(a);
    
```

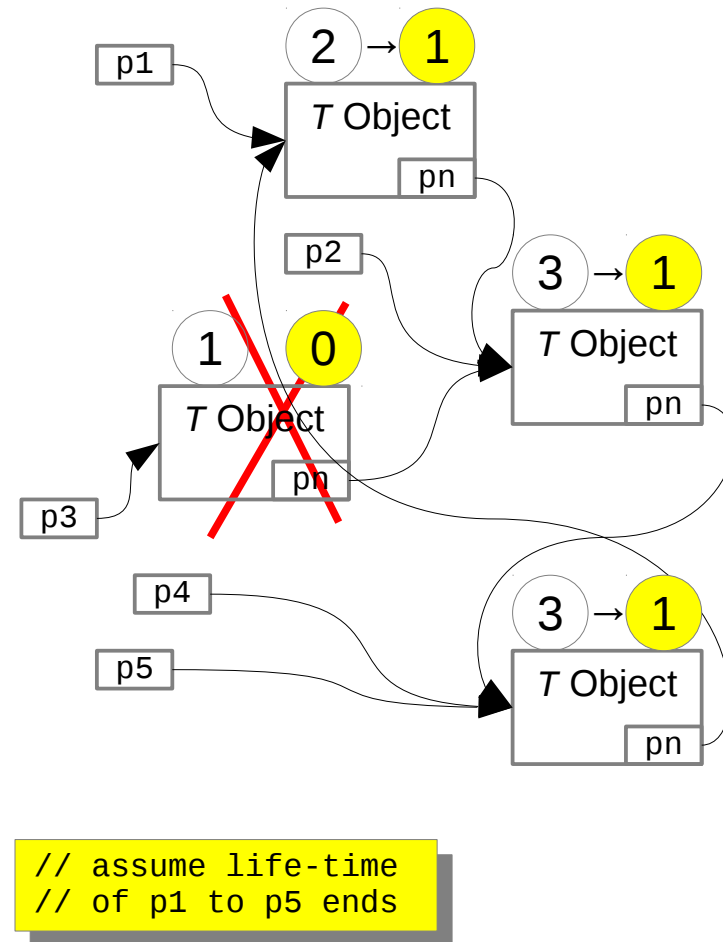
generally compiles but will slice Apple-s and Banana-s to their Fruit base

OK but, just a Fruit ... (Sorry Sir, doesn't taste like a Banana anymore)

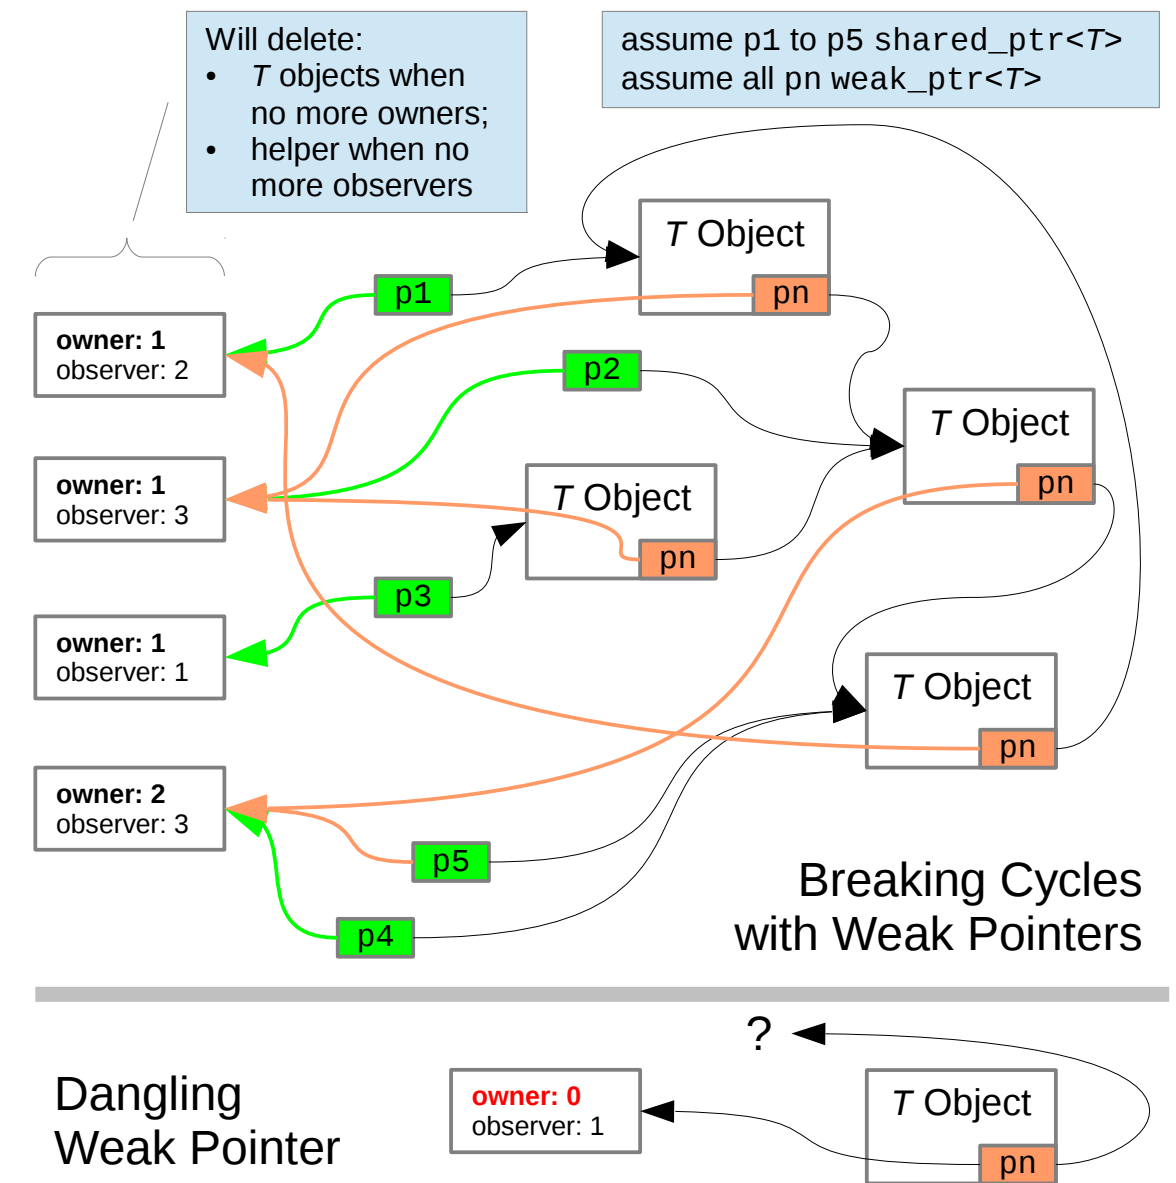
generally compiles but probably does not what is expected because only the Fruit-part of the apple gets modified



Reference Counting Principle



Problem of Cyclic References



Comparing ...	<code>std::unique_ptr<T></code>	<code>std::shared_ptr<T></code>	Remarks
Characteristic	refers to a single object of type <code>T</code> , uniquely owned	refers to a single object of type <code>T</code> , possibly shared with other referrers	may also refer to "no object" (like a <code>nullptr</code>)
Data Size	same as plain pointer	same as a plain pointer plus some extra space per referred-to object	
Copy Constructor	no*	yes	particularly efficient as only pointers are involved
Move Constructor	yes		
Copy Assignment	no*		a <code>T</code> destructor must also be called in an assignment if the current referrer is the only one referring to the object
Move Assignment	yes		
Destructor (when referrer life-time ends)	always called for referred-to object	called for referred-to object when referrer is the last (and only) one	

*: explicit use of `std::move` for argument is possible

Smart Pointer Comparison

(cc) BY-SA: Technische Beratung für EDV, Dipl.-Ing. Martin Weitzel, Germany, <http://tbfe.de>

